

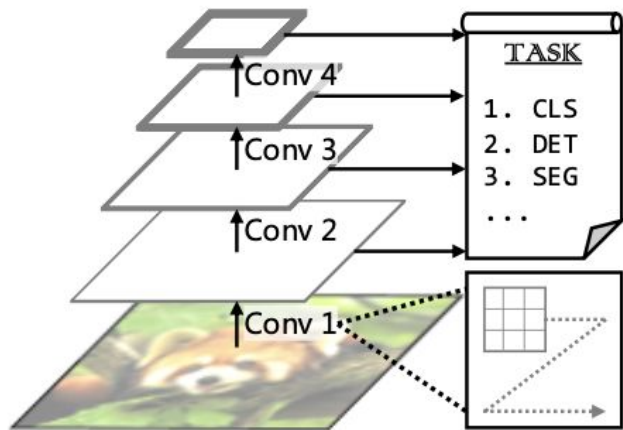
Bottleneck Transformers for Visual Recognition

A Self-attention model for vision

Will, Tin, Tarunyaa

Premise: Out with CNNs


Better alternatives to pure CNNs for CV such as object detection, image classification



(a) CNNs: VGG [53], ResNet [21], etc.

→ ResNet: Residual Neural Network

Intermediate input added to the output of a series of convolutional blocks to enable scaling.



They capture local patterns information but they **fail to understand long-term dependencies + require many layers.**

Premise: Out with CNNs

Self-attention to model long-term dependencies

- Longer-term memory than RNNs & LSTMs
- NLP to Image processing
- GPT, BERT

They capture long term dependencies and don't require as many layers.

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .

The FBI is chasing a criminal on the run .


The FBI is chasing a criminal on the run .

Proposal: Using Self-Attention in Vision


HYBRID SOLUTION: Replace spatial convolutional layers w/ multi-head self attention layer

- Use convolutions to deal with large images efficiently; do spatial downsampling
- Letting global self-attention work on small resolutions

Avoids processing large images w/ self-attention since its memory & computation required scales quadratically w/ spatial dimensions



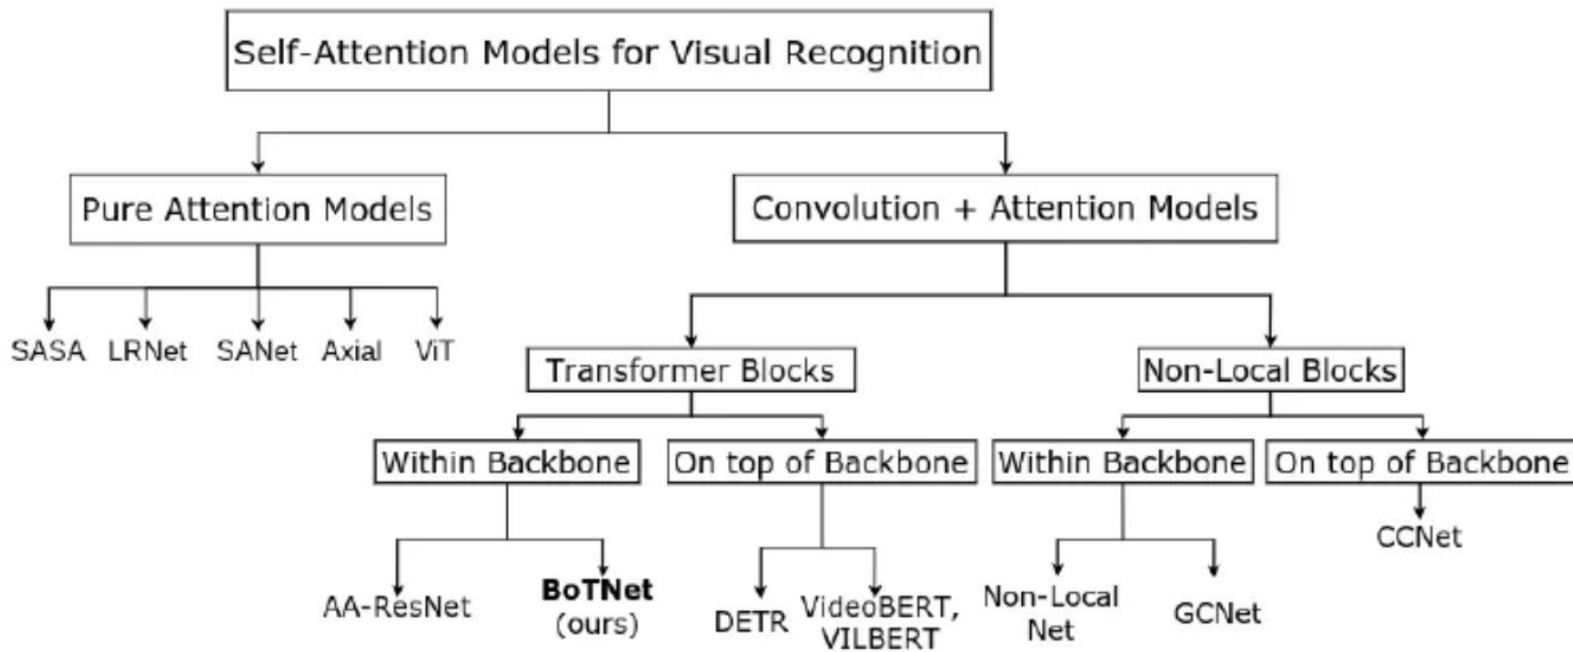
ResNet Bottleneck



Bottleneck Transformer

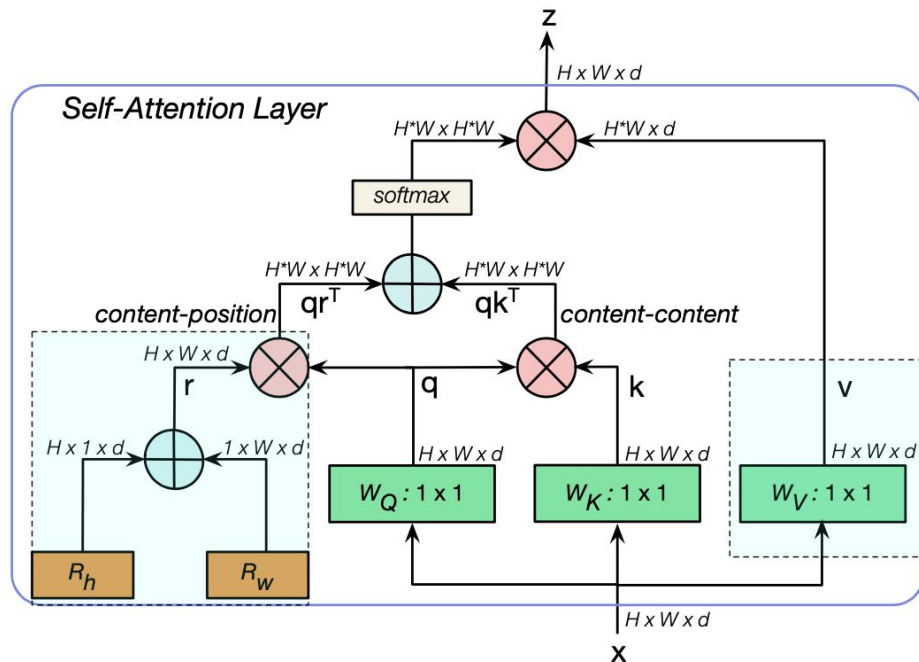
Proposal: Using Self-Attention in Vision

HYBRID SOLUTION: Replace spatial convolutional layers w/ multi-head self attention layer



Method: Positional Encoding

Making the attention-operation *position aware*



- Global attention is performed on a 2D feature map
- Split relative position encodings, R_h and R_w , for height and width respectively

Method: Model Architecture

Low relative overhead

stage	output	ResNet-50	BoTNet-50
c1	512×512	$7 \times 7, 64, \text{stride } 2$	$7 \times 7, 64, \text{stride } 2$
c2	256×256	$3 \times 3 \text{ max pool, stride } 2$	$3 \times 3 \text{ max pool, stride } 2$
		$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
c3	128×128	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$
c4	64×64	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$
c5	32×32	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ \text{MHSA}, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
# params.		25.5×10^6	20.8×10^6
M.Adds		85.4×10^9	102.98×10^9
TPU steptime		786.5 ms	1032.66 ms

→ Only difference is the use MHSA layer in c5

→ BoT50 has only 1.2 x multiple-adds and 1.3 x training overheads with 1.2x fewer parameters.

Results: Comparison w/ ResNet on Coco

*BoT50 is better than R50 and R101,
competitive with R152*

Backbone	AP ^{bb}	AP ^{mk}
R50	42.1	37.7
BoT50	43.6 (+ 1.5)	38.9 (+ 1.2)
R101	43.3	38.4
BoT101	45.5 (+ 2.2)	40.4 (+ 2.0)
R152	44.2	39.1
BoT152	46.0 (+ 1.8)	40.6 (+ 1.5)

*Relative positional encoding
boosts performance*


Backbone	Att. Type	AP ^{bb}	AP ^{mk}
R50	-	42.1	37.7
BoT50	qk^T	42.7 (+ 0.6)	38.3 (+ 0.6)
BoT50	qr_{relative}^T	43.1 (+ 1.0)	38.4 (+ 0.7)
BoT50	$qk^T + qr_{\text{relative}}^T$	43.6 (+ 1.5)	38.9 (+ 1.2)
BoT50	$qk^T + qr_{\text{abs}}^T$	42.5 (+ 0.4)	38.1 (+ 0.4)

Surpasses previous best published model on ResNet

Results: Comparison w/ ResNet on Coco

BoTNet benefits from training on larger images

Backbone	res	AP ^{bb}	AP ^{mk}
R50	1280	44.0	39.5
BoT50	1024	45.9 (+ 1.9)	40.7 (+ 1.2)
BoT50	1280	46.1 (+ 2.1)	41.2 (+ 1.8)
R101	1280	46.4	41.2
BoT101	1024	47.4 (+ 1.0)	42.0 (+ 0.8)
BoT101	1280	47.9 (+ 1.5)	42.4 (+ 1.2)

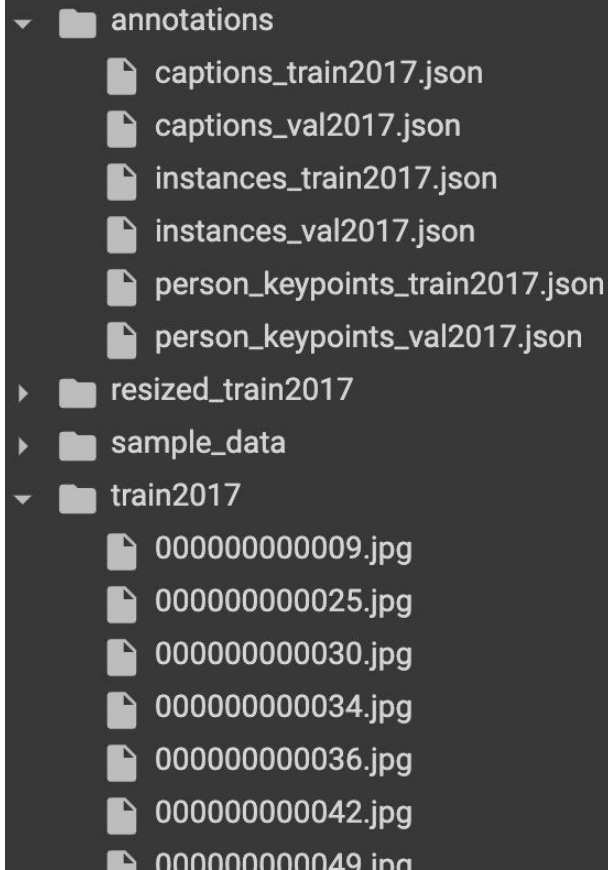


Scales well w/ larger images

COCO dataset

```
annotation{
  "id"           : int,
  "image_id"     : int,
  "category_id"  : int,
  "segmentation" : RLE or [polygon],
  "area"         : float,
  "bbox"         : [x,y,width,height],
  "iscrowd"      : 0 or 1,
}


categories[{
  "id"           : int,
  "name"         : str,
  "supercategory": str,
}]
```



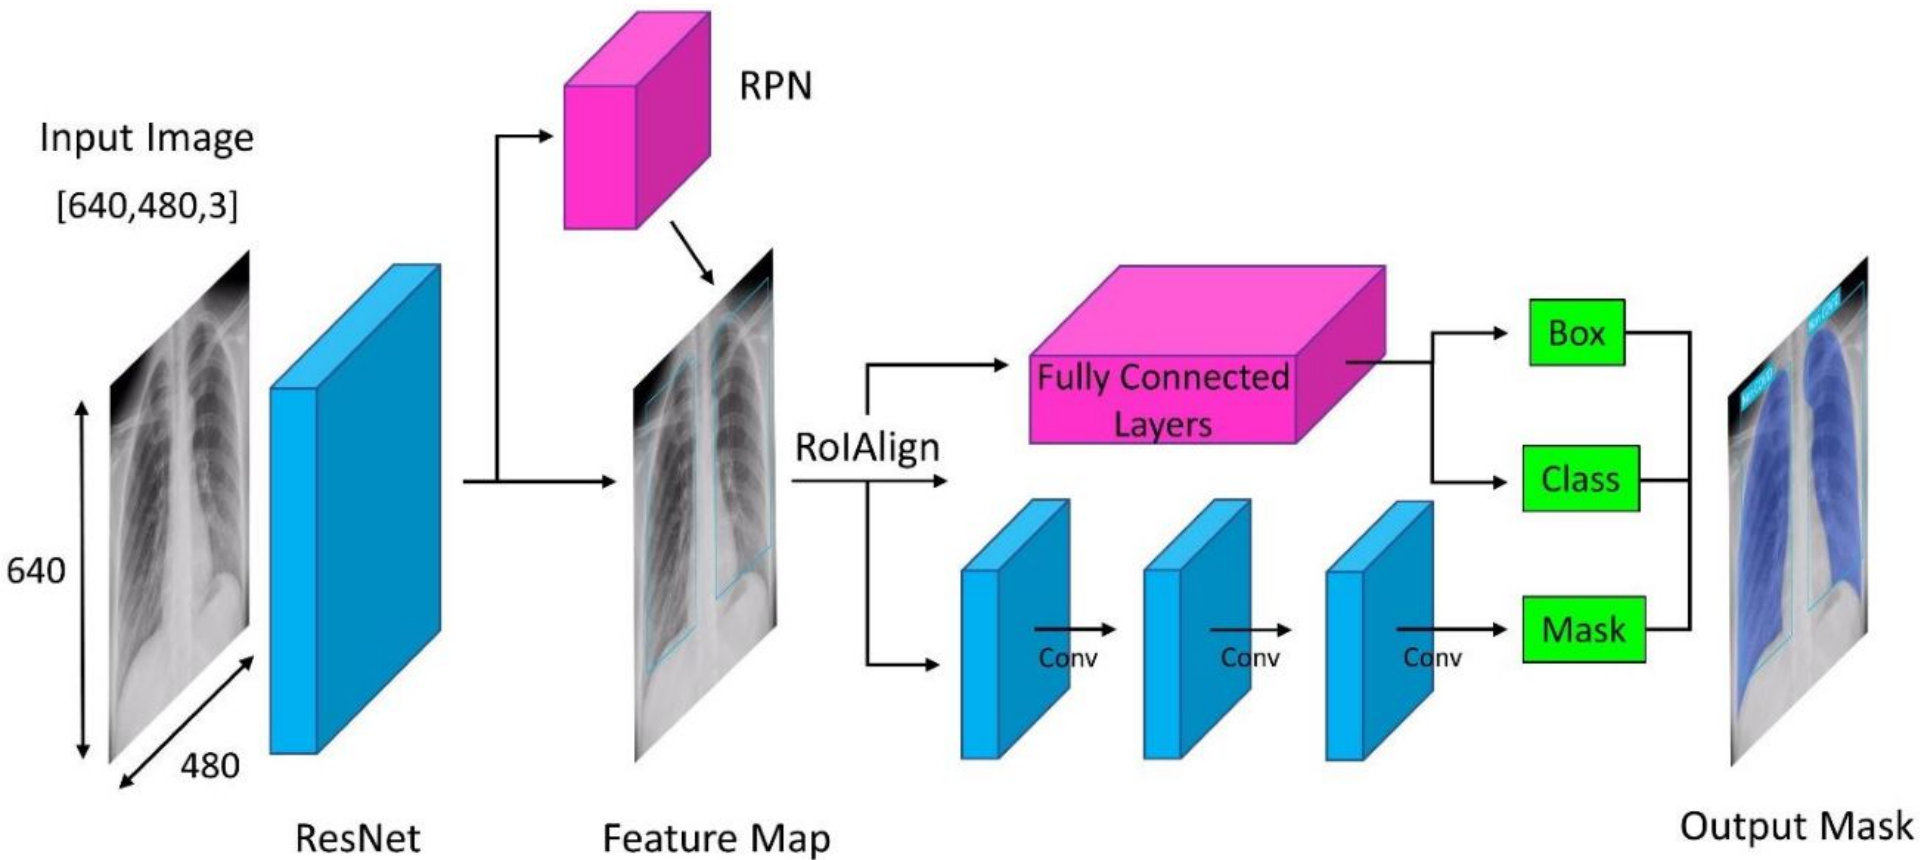
- ▼ annotations
 - captions_train2017.json
 - captions_val2017.json
 - instances_train2017.json
 - instances_val2017.json
 - person_keypoints_train2017.json
 - person_keypoints_val2017.json
- ▶ resized_train2017
- ▶ sample_data
- ▼ train2017
 - 000000000009.jpg
 - 000000000025.jpg
 - 000000000030.jpg
 - 000000000034.jpg
 - 000000000036.jpg
 - 000000000042.jpg
 - 000000000049.jpg



Instance segmentation



Object detection



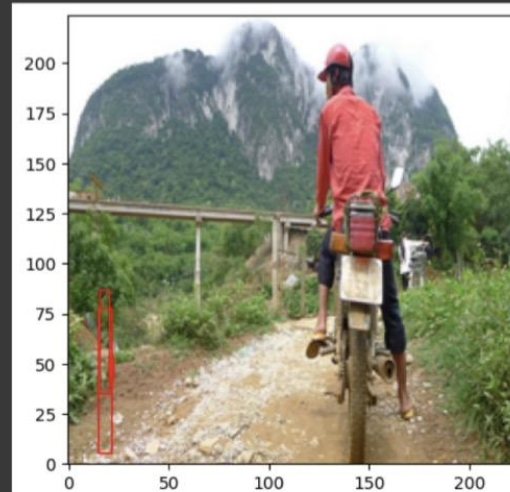
Mask R-CNN Architecture

MM Detection



```
1 import torch.nn as nn
2 from torchvision.models.detection import MaskRCNN
3 from torchvision.models.detection.rpn import AnchorGenerator
4 from torchvision.ops import MultiScaleRoIAlign
5
6
7 backbone = ResNet(Bottleneck, [3, 4, 6, 3], resolution=(224, 224), heads=4)
8
9 # you are effectively informing the rest of the MaskRCNN model about the shape
10 # of the tensors it will receive from the backbone. This ensures that subsequent
11 # layers can be correctly configured to work with these tensors.
12 backbone.out_channels = 2048
13
14 anchor_generator = AnchorGenerator(sizes=((32, 64, 128, 256, 512,)), aspect_ratios=((0.5, 1.0, 2.0,)))
15
16 roi_pooler = MultiScaleRoIAlign(featmap_names='0', output_size=7, sampling_ratio=2)
17
18 mask_roi_pooler = MultiScaleRoIAlign(featmap_names='0', output_size=14, sampling_ratio=2)
19
20 # Define the model
21 model = MaskRCNN(backbone, num_classes=91, # COCO has 80 classes + background
22                 rpn_anchor_generator=anchor_generator,
23                 box_roi_pool=roi_pooler,
24                 mask_roi_pool=mask_roi_pooler)
```

(2, 224, 224, 3)



Our Code

- Image Classification on CIFAR
- Does slightly better than CNN and uses fewer parameters
- Exploration of the architecture
- [Here](#)

Thank you!